

# Concurrent Transactions

# Concurrent Transactions

## Motivation

- The main objective of developing database is to enable many users share data concurrently.
- Concurrent access is relatively easy if all users are only reading data; however when users are updating data, there may be interference (conflicts) that can result to inconsistencies

# Concurrent Transactions

## Definitions:

### Concurrency

- In computer science, concurrency is a property of systems in which several computations are executing simultaneously and potentially interacting with each other.

### Concurrency control

- It's the process of managing simultaneous operations on the data base without having the interfere with each another.

# Concurrent Transactions

## Purpose of Concurrency Control

- To Enforce isolation among conflicting transactions
- To Preserve database consistency
- To resolve conflicts among conflicting transactions

# Concurrent Transactions

## Conflicting Operations

- Two operations on the same data item in different transactions conflict if one of them is a write.
- If we denote database read and write operations on a data item  $x$  as  $r(x)$  and  $w(x)$ , then we say the operations conflict:
  1. If the operations belong to different transactions
    - One transaction can not conflict
  2. If the operations must access the same data item and
  3. At least one of the operation should be write

# Concurrent Transactions

## Conflicting Operations

- Two operations on the same data item conflict if at least one of the operations is a write:
  - $r(x)$  and  $w(x)$  conflict
  - $w(x)$  and  $r(x)$  conflict
  - $w(x)$  and  $w(x)$  conflict
  - $r(x)$  and  $r(x)$  do not
  - $r/w(x)$  and  $r/w(y)$  do not
- Order of conflicting operations matters
  - If  $T1.r(a)$  precedes  $T2.w(a)$ , then conceptually,  $T1$  should precede  $T2$

# Concurrent Transactions

## Conflicting Operations~example


Consider these two different sequences of execution :

- T1

$A = A + 100$   
 $B = B - 100$   
COMMIT

- T2

$A = A * 1.05$   
 $B = B * 1.05$   
COMMIT



---

**A = 210 DH**

**B = 0 DH**

---

**A = 210 DH**

**B = 5 DH**

# Concurrent Transactions

## Examples of Interference Between Concurrent Transactions

- There are many ways in which concurrently executing transactions can interfere with one another and so compromise the integrity and consistency of the DB.
- Three examples of such interference are:
  - Lost Update Problem
  - Uncommitted Dependency Problem
  - Inconsistency Analysis Problem



# Concurrent Transactions

## The Lost Update Problem

- An apparently successfully completed update operation by one user can be overridden by another user. This is known as the Lost Update Problem.

# Concurrent Transactions

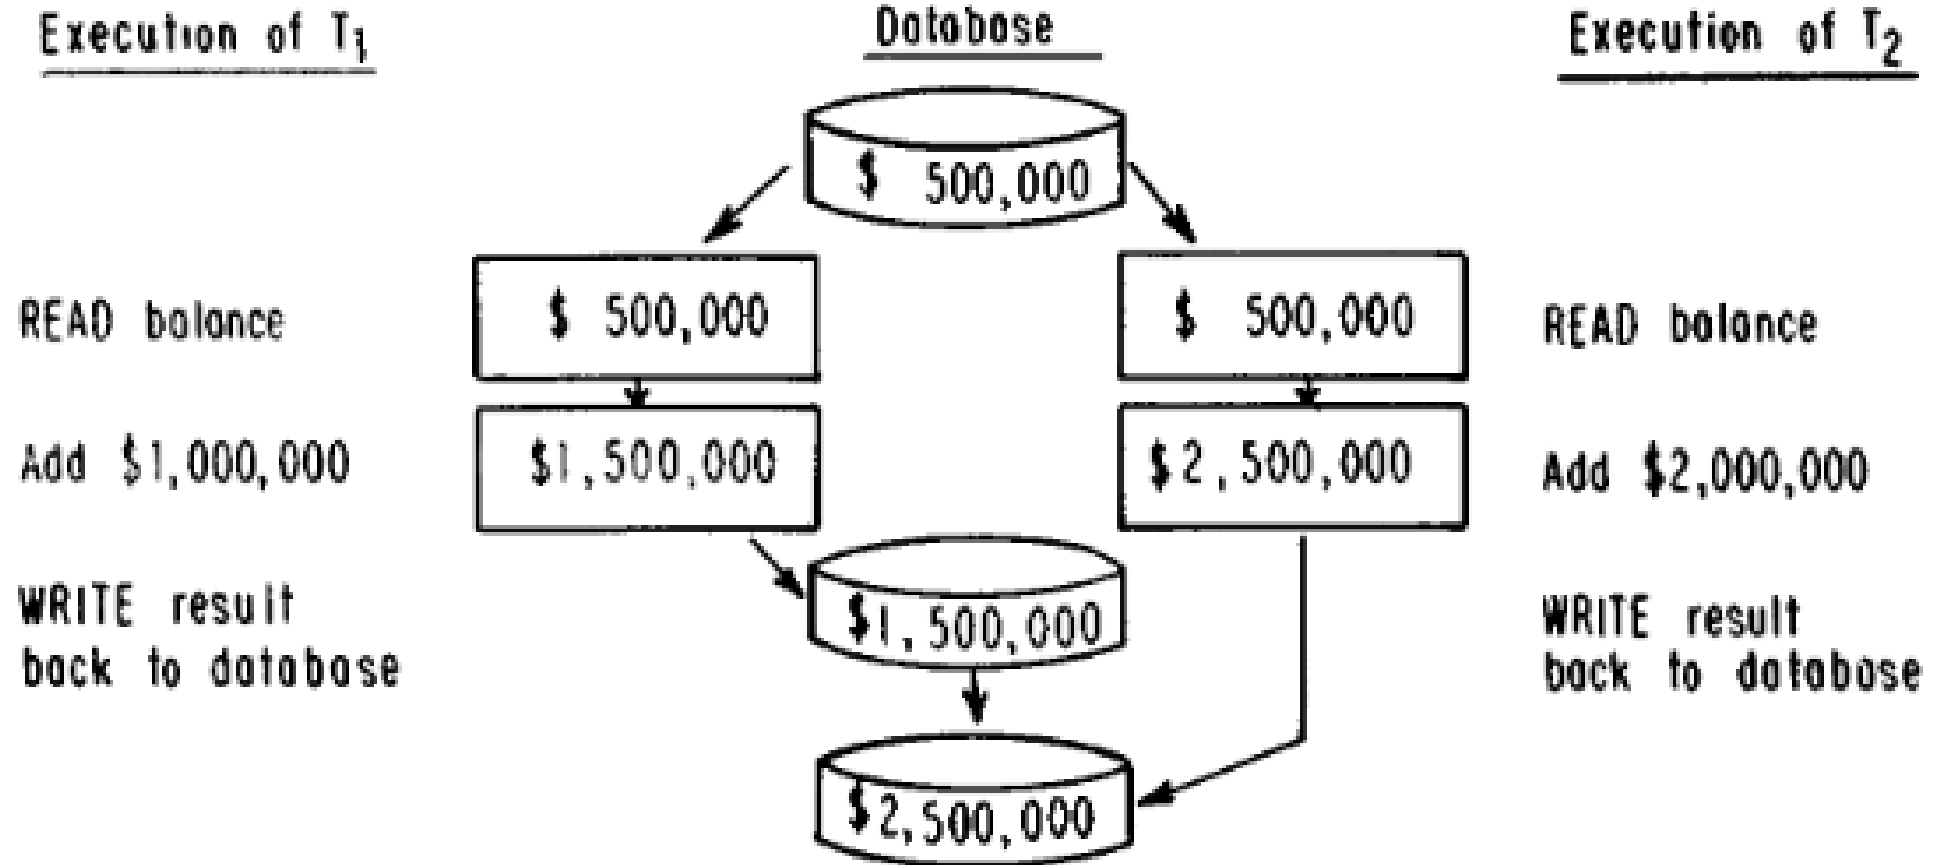
## The Lost Update Problem

**Example I:** Suppose two customers simultaneously try to deposit money into the same account. The two ATMs handling the two customers could read the account balance at approximately the same time, compute new balances in parallel, and then store the new balances back into the database.

***Take Note that:*** In response to customer requests, ATMs retrieve data from a database, perform computations, and store results back into the database.

# Concurrent Transactions

## The Lost Update Problem



# Concurrent Transactions

## The Lost Update Problem

### Example I

<u>Time</u>	<u>Transaction A</u>	<u>Transaction B</u>
t1	Read (R)	~
t2	~	Read (R)
t3	Write (R)	~
t4	~	Write (R)

- Transaction A's update is lost at time t4 because transaction B overwrites it without even looking at it!

# Concurrent Transactions

## The Lost Update Problem

*Example II:* Transaction  $T_1$  is executed concurrently with transaction  $T_2$ .

Time	$T_1$	$T_2$	Bal <sub>x</sub>
$t_1$	-	begin_transaction	100
$t_2$	begin_transaction	read(bal <sub>x</sub> )	100
$t_3$	read(bal <sub>x</sub> )	bal <sub>x</sub> = bal <sub>x</sub> + 100	100
$t_4$	bal <sub>x</sub> = bal <sub>x</sub> - 10	write(bal <sub>x</sub> )	200
$t_5$	write(bal <sub>x</sub> )	Commit	90
$t_6$	Commit	-	90

- Transaction  $T_1$  updates the same record updated earlier by  $T_2$  at time  $t_5$  on the basis of values read at  $t_3$ .
- Question: How do we solve the lost update problem?

# Concurrent Transactions

## Uncommitted Dependency Problem

- A violation of the integrity constraints governing the database can arise when two transactions are allowed to execute concurrently without being synchronized.
- Uncommitted dependency Problem occurs when one transaction is allowed to see the intermediate results of another transaction before it has committed.

# Concurrent Transactions

## Uncommitted Dependency Problem

**Example I:** Consider this situation:

Time	T <sub>3</sub>	T <sub>4</sub>	Bal <sub>x</sub>
t <sub>1</sub>	-	begin_transaction	100
t <sub>2</sub>	-	read(bal <sub>x</sub> )	100
t <sub>3</sub>	-	bal <sub>x</sub> = bal <sub>x</sub> + 100	100
t <sub>4</sub>	begin_transaction	write(bal <sub>x</sub> )	200
t <sub>5</sub>	read(bal <sub>x</sub> )	:	200
t <sub>6</sub>	bal <sub>x</sub> = bal <sub>x</sub> - 10	rollback	100
t <sub>7</sub>	write(bal <sub>x</sub> )	-	190
t <sub>8</sub>	commit	-	190

**Question:** How do we solve the uncommitted dependency problem?

# Concurrent Transactions

## Uncommitted Dependency Problem

- Transaction  $T_3$  reads on uncommitted update at time  $t_5$ .
- The update is then undone at time  $t_6$
- Transaction  $T_3$  is therefore operating on false assumption i.e. Transaction  $T_3$  becomes dependent on an uncommitted update at time  $T_2$



# Concurrent Transactions

## Inconsistent Analysis Problem

- Transactions that only read the database can obtain inaccurate results if they are allowed to read partial results of incomplete transactions, which are simultaneously updating the database
- Problem of inconsistent analysis occurs when a transaction reads several values from the database but a second transaction updates some of them during the execution of the first.

# Concurrent Transactions

## Inconsistent Analysis Problem

- Consider two transactions A and B operating on bank account records
- Transaction B is summing account balances.
- Transaction A is transferring an amount 10 from account 1 to account 3
- Account 1 = 100; Account 2 = 50; Account 3 = 25

# Concurrent Transactions

## Inconsistent Analysis Problem

Time	Transaction A	Transaction B	Acc1	Acc2	Acc3	Sum
t <sub>1</sub>	Read Acc1	Read Acc1	100	50	25	0
t <sub>2</sub>	Acc1 ~ 10	Sum + Acc1	100	50	25	100
t <sub>3</sub>	Write Acc1	Read Acc2	90	50	25	100
t <sub>4</sub>	Read Acc3	Sum + Acc2	90	50	25	150
t <sub>5</sub>	Acc3 + 10	~	90	50	25	150
t <sub>6</sub>	Write Acc3	~	90	50	35	150
t <sub>6</sub>	Commit	Read Acc3	90	50	35	150
t <sub>7</sub>	~	Sum + Acc3	90	50	35	185
t <sub>8</sub>	~	Commit	90	50	35	185

✉ Sum = 185 not 175!!

Question: How do we solve the inconsistent analysis problem?

# How to solve Interference Between Concurrent Transactions

- The objective of concurrency control protocol is to schedule transactions in such a way as to avoid any interference between them.
- When there are multiple transactions that are running in a concurrent manner and the order of operation is needed to be set so that the operations do not overlap each other, Scheduling is brought into play and the transactions are timed accordingly.
  - A quick solution is to allow only one transaction to be executed and committed before another transaction begins to execute.
  - However, the aim of multi-user DBMS is to maximize the degree of concurrency or parallelism in the system without interference among the transactions.



# References

- <https://www.geeksforgeeks.org/types-of-schedules-in-dbms/>
- <http://www.ccs.neu.edu/home/kathleen/classes/cs3200/9-Transactions.pdf>